

Privilege separation

Modern computers are


1. multi-users
2. multi-tasking

Goal: Prevent potentially misbehaving users and/or applications from harming the rest of the system

Permissions system: mechanisms for achieving separation between components

Key assumptions for separation

1. The system know who the user is - user has authenticated, e.g. using username / password
2. Complete mediation - all requests are mediated - all requests go to the reference monitor that enforces specified access control policies



The **reference monitor** grants permission to **users** to apply certain **operations** to a given **resource**

Central question

"Who is allowed to access **what** and **how?**"

The subject (who) - eg. user, application, process

The object (what) - protected resource, eg. hardware device, network socket, memory, files, directories,

The access operation (how) - eg. read, write, execute

Users

Two types of accounts each with a unique identifier, the user ID (uid):

1. User accounts - associated with humans
2. Service accounts - associated with background processes

```

narapint@nyrto-thinkpad:~$ more /etc/passwd
root:x:0:0:root:/root:/bin/bash
daemon:x:1:1:daemon:/usr/sbin:/usr/sbin/nologin
bin:x:2:2:bin:/bin:/usr/sbin/nologin
sys:x:3:3:sys:/dev:/usr/sbin/nologin
sync:x:4:65534:sync:/bin:/bin/sync
games:x:5:60:games:/usr/games:/usr/sbin/nologin
man:x:6:12:man:/var/cache/man:/usr/sbin/nologin
lp:x:7:7:lp:/var/spool/lpd:/usr/sbin/nologin
mail:x:8:8:mail:/var/mail:/usr/sbin/nologin
news:x:9:9:news:/var/spool/news:/usr/sbin/nologin
uucp:x:10:10:uucp:/var/spool/uucp:/usr/sbin/nologin
proxy:x:13:13:proxy:/bin:/usr/sbin/nologin
  
```

— You can list ALL accounts in a UNIX machine by inspecting the `passwd` file that can be found in `/etc/passwd`

- ▶ One entry in the `/etc/passwd` per account with the fields: `username:password:uid:gid:uid_info:home:shell`
- ▶ uid 0 - user root uid group id

Groups

- Groups are sets of users that share resources
- Every group has a name and a unique identifier, the group ID (gid)
- Allow for easier users management and monitoring

You can list ALL groups by inspecting the grp. file in /etc/group

```
narapin@myrto-thinkpad:~$ more /etc/group
root:x:0:
daemon:x:1:
bin:x:2:
sys:x:3:
adm:x:4:syslog,narapin
tty:x:5:
disk:x:6:
lp:x:7:
mail:x:8:
```

There are 2 users: syslog and 'myself'

- One entry in the /etc/group per group with the fields:

group_name:password:gid:group_list

When you specify a directory, you can only specify one group

Directory

(slide) pdf

Directory permissions

- Execute permission on a directory allows traversing it
- Read permission on a directory allows lookup

2.9. Quizz: Imagine you have the following groups:

- infr10067 - for any user involved with the Computer Security course
- tas - for all Informatics TAs

TAs on INFR10067

How can you have a folder only for Computer Security TAs?

```
narapin@myrto-thinkpad:~/Documents/Work/Teaching/INFR10067-ComputerSecurity/2021/Lectures/L18.AccessControl/conjunction$ ls -l
total 4
drwxr-xr-x 3 narapin tas 4096 Feb 23 22:50 only_for_tas
narapin@myrto-thinkpad:~/Documents/Work/Teaching/INFR10067-ComputerSecurity/2021/Lectures/L18.AccessControl/conjunction$ ls -l only_for_tas/
total 4
drwxr-xr-x 2 narapin infr10067 4096 Feb 23 22:50 only_for_infr10067_tas
narapin@myrto-thinkpad:~/Documents/Work/Teaching/INFR10067-ComputerSecurity/2021/Lectures/L18.AccessControl/conjunction$
```

Nest 2 directories! 2 directories: only_for_tas & only_for_infr10067_tas

In only_for_tas, I've given execute permission only to the grp tas. Inside!

In only_for_infr10067_tas, I've given execute permission only to the users in infr10067

File permissions

- All resources (sockets, directories, files) are managed as files
- 3 defined permissions read (r), write (w), execute (x)
- Permissions are defined for the owner, the owner's group, and other users
- Root and owner can change file permissions
- Only root can change file ownership

```
narapin@myrto-thinkpad:~/Documents/Work/Teaching/INFR10067-ComputerSecurity/2021/Lectures/L18.AccessControl$ ls -l
total 352
drwxr-xr-x 2 narapin narapin 4096 Feb 23 17:27 Images
-rw-r--r-- 1 narapin narapin 1839 Feb 23 17:31 L18.AccessControl.aux
-rw-r--r-- 1 narapin narapin 47121 Feb 23 17:31 L18.AccessControl.log
-rw-r--r-- 1 narapin narapin 835 Feb 23 17:31 L18.AccessControl.nav
-rw-r--r-- 1 narapin narapin 0 Feb 23 17:31 L18.AccessControl.out
-rw-r--r-- 1 narapin narapin 258111 Feb 23 17:31 L18.AccessControl.pdf
-rw-r--r-- 1 narapin narapin 0 Feb 23 17:31 L18.AccessControl.snm
-rw-r--r-- 1 narapin narapin 9769 Feb 23 18:05 L18.AccessControl.tex
-rw-r--r-- 1 narapin narapin 23638 Feb 20 01:28 L18.AccessControl.tex
-rw-r--r-- 1 narapin narapin 0 Feb 23 17:31 L18.AccessControl.tex
```

The owner (myself) has read & write permissions, NO execute permission

The rest of the group ONLY has read permission

The other users ONLY have read permission

Processes

- Each process has a unique identifier, the process ID (pid)
- Each process is associated with the user that spanned it

The init process which is run by the root user is given pid of 0

```
narapin@myrto-thinkpad:~$ ps -ef
UID          PID     PPID  C  STIME TTY          TIME CMD
root           1         0  0   Feb22 ?        00:00:43 /sbin/init splash
root           2         0  0   Feb22 ?        00:00:00 [kthreadd]
root           4         2  0   Feb22 ?        00:00:00 [kworker/0:0H]
root           6         2  0   Feb22 ?        00:00:00 [rm_percpu_wq]
root           7         2  0   Feb22 ?        00:00:00 [ksoftirqd/0]
root           8         2  0   Feb22 ?        00:00:00 [rcu_sched]
root           9         2  0   Feb22 ?        00:00:00 [rcu_bh]
root          10         2  0   Feb22 ?        00:00:00 [migration/0]
root          11         2  0   Feb22 ?        00:00:00 [watchdog/0]
root          12         2  0   Feb22 ?        00:00:00 [cpuhp/0]
root          13         2  0   Feb22 ?        00:00:00 [cpuhp/1]
root          14         2  0   Feb22 ?        00:00:00 [watchdog/1]
```

- When a user runs a process, it runs with that user's privileges, i.e. they can access any resource that user has permissions for
- By default, a child process inherits its parent's privileges
- Processes are isolated in memory

Process user IDs

Every process has:

- ▶ Real user ID (uid) - the user ID that started that process
- ▶ Effective user ID (euid) - the user ID that determines the process privileges
- ▶ Saved user ID (suid) - the effective user ID before the last modification (It is possible to change the euid)

Users can change a process' IDs:

```
setuid(x)
uid ← x
euid ← x
suid ← x
```

```
seteuid(x)
uid ← uid
[euid ← x]
suid ← suid
```

Only the euid changes!

- ▶ Root can change euid/uid to arbitrary values x:
- ▶ Unprivileged users can only change euid to uid or suid:

only drop the privileges

10/15

Elevating privileges - setuid programs

- ▶ An executable file can have the set-user-ID property (setuid) enabled
- ▶ If A executes a setuid file owned by B, then the euid of the process is B and not A
- ▶ Writing secure setuid programs is tricky because vulnerabilities may be exploited by malicious user actions
- ▶ Some programs that access system resources are owned by root and have the setuid bit set (setuid programs)

```
serapi@myrta-thinkpad:~/bin$ ls -l | grep passwd
-rwxr-xr-x 1 root root 75804 Jan 25 2018 passwd
-rwxr-xr-x 1 root root 749976 Feb 7 23:28 grub-kpseed pkdfz
-rwxr-xr-x 1 root root 39648 Jan 25 2018 passwd
serapi@myrta-thinkpad:~/bin$
```

```
serapi@myrta-thinkpad:~/bin$ ls -l | grep shadow
-rwxr-xr-x 1 root shadow 860 Feb 27 18:11 gshadow
-rwxr-xr-x 1 root shadow 845 Sep 21 15:00 gshadow
-rwxr-xr-x 1 root shadow 1172 Feb 27 18:11 shadow
-rwxr-xr-x 1 root shadow 1172 Feb 27 18:11 shadow
serapi@myrta-thinkpad:~/bin$
```

If setuid bit is set, tht. prog. runs w/ the euid of its owner, rather than the process tht. executed it.

Processes are created by a mechanism called forking, where a new process is created (that is, forked) by an existing process.

Dropping privileges with setuid

Imagine a program that runs as root and wants to fork a process with lower privileges using the following code:

```
if (auth(uid, pwd) == SUCCESS) {
    if (fork() == 0) {
        setuid(uid);
        exec("/bin/bash");
    }
}
```

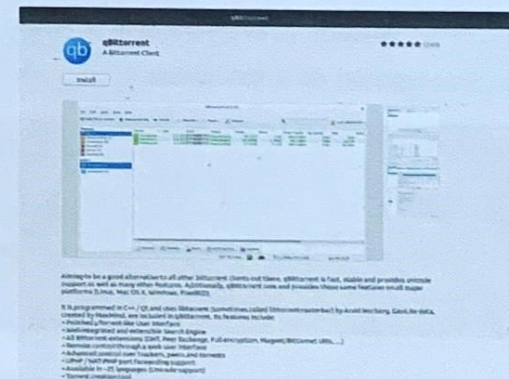
This is possible bc. after seteuid(uid), the suid does not change! Hence, unprivileged user can change euid to suid and become root!!

What we can do instead is:
setuid(uid);

UNIX permissions are too coarse-grained

Mobile OS (i.e. Android) does things differently

All application installed by a single user account have the same privileges!



I have given to this program all my privileges

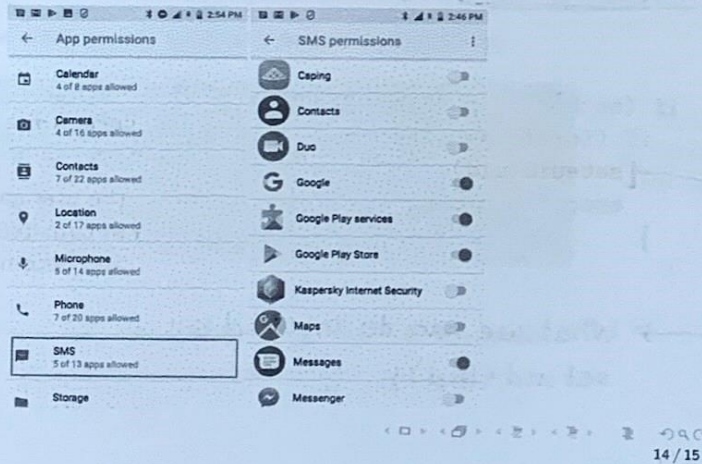
!!? What if qBittorrent is malware ?!!

- ▶ Better delegate capabilities associated with specific root powers

13/15

Android permissions

- ▶ Each app runs with a different user ID
- ▶ Apps do not interact
- ▶ Permissions are set per app



As long as the OS is not compromised (the attacker does NOT have root privileges), then installing a malicious app should NOT affect other apps.

Take aways

The UNIX security model provides a simple and flexible model, but permissions are too coarse-grained:

- same permissions for all applications ran under a single user account

- many utilities have the setuid bit enabled

These setuid progs. are opportunities for an attacker to gain another user's privileges (i.e. root privileges)

→ many opportunities for privilege escalation attacks

→ better use capabilities when delegating privileges